Module 1: Slide Deck Video Transcription

1

Hi, I'm Jordan Reyes, and I lead AI training here at Tallgrass Advisory. A quick note before we start: I'm a fictional character presented by an AI avatar. That's on purpose. This whole course is about using AI well and being open about it. Velox Motorworks has hired Tallgrass to get its new finance team ready for AI-first work, and that's you. For the next three weeks, you're with me.

2

Velox's CFO, Priya Natarajan, has a big goal. A competitor just cut its month-end close in half using AI, and Priya wants Velox to lead, not follow. She's asking every new hire to approach every task with one question: how can AI help me do this? That question is the habit we'll build all term. People who ask it get more done, and they get noticed.

3

So what are you working with? Tools like Claude are large language models. They learned patterns from enormous amounts of text, and they generate answers by predicting what should come next. That's why they write, summarize, and explain so well. It's also why they can sound completely confident while being wrong. Fluent isn't the same as correct.

4

Here's where AI shines: drafting emails and memos, summarizing long documents, explaining a concept you've never seen, translating, organizing messy notes, and taking a first pass at data you hand it. In finance, that's hours back every week.

5

And here's where you check it. Large calculations. Current facts, like this year's tax rate. Citations, which it can invent. And anything it simply can't see, like Velox's own records. When AI fills a gap with something made up, we call that a hallucination. Champions don't avoid AI because of this. They verify. Great answer? Now let's prove it.

6

This week you'll sort 10 real Velox tasks into three categories. AI does it well: low risk, a quick read is enough. AI helps, a person checks: AI does the work, you verify the facts and numbers. And AI assists, a person decides: AI prepares the analysis, but a named

person makes the call or signs. Every task gets AI help. What changes is how much you check and who owns the decision.

7

Let me show you how I'd sort one, live. This isn't one of your 10 tasks, so no spoilers. Walt Brenner, our Controller, wants a summary of every travel expense over $5,000 from last month. I've got a simulated export of 300 expense lines. First, the prompt. I won't just type "summarize this." I'll give context: "You're helping the Velox Controller review travel spending. From the attached file, list every expense over $5,000 with the employee's department, the amount, and the purpose. Then total them by department, and flag anything that looks unusual." Notice that I named the audience, the threshold, the fields I want, and what to flag.

8

Here's the output. It's fast and well organized: a list of the 14 expenses over $5,000, totals by department, and a few flags. That's genuinely useful. Ten minutes of work done in seconds. Now let's prove it. I'll filter the original file myself for amounts over $5,000. I count 14, the same as Claude. I'll add them up: $88,070 in total, and the department totals match too. So the numbers hold up. Then I read Claude's comments, and one stands out. It points out two expenses near the line: one for $5,025, which is on the list, and one for exactly $5,000.00, which isn't, because my prompt said over $5,000, not at or over. Claude did exactly what I asked, and it told me about the edge case. Good. But look at who has to decide: should an expense of exactly $5,000 count? That depends on Velox's approval limit, and that's Walt's call, not Claude's. My wording decided the answer.

9

One more thing: the file had employee names in it. Because this is an approved internal tool and the data is simulated, that's fine here. With real data in an unapproved tool, I'd strip the names first. So where does this task go? AI helps, a person checks. AI does the heavy lifting; I verify the count and the totals. Walt decides what to do about the flagged items, and whether the $5,000 line includes $5,000. That's the kind of reasoning I want to see on all 10 of your tasks: what AI did well, what you checked, and who owns the decision.

10

Velox has five rules for AI. Keep confidential data out of AI tools unless they're approved, and anonymize first. Check accuracy. Watch for bias. Disclose when AI helped. And remember that a person is always accountable. One of your 10 tasks involves employee salary data, and how you handle it matters as much as the answer you get.

11

Here's your week. You'll try AI on a task before I teach you anything. Then we'll learn, work through ethics, and you'll build your own coach: a version of me in a Claude Project that gives you feedback before you submit. You'll revise, submit your best work, take a short quiz, and reflect. At the end, you'll hand Priya a memo she can act on. Let's get to work.