Module 3 - Slide Deck Video Transcription

1 Hi, I'm Jordan Reyes from Tallgrass Advisory. Quick reminder: I'm a fictional character presented by an AI avatar, on purpose, because this course is about using AI well and being open about it. Last week you learned to give AI the right context. This week the job gets bigger. Walt Brenner, Velox's Controller, wants help with a task that has many steps, and he's worried that a wrong answer could look right. Your job is to show him a way to use AI that lets him see how it got there.

2 Here's Walt's problem. Every month his staff accountant reconciles Velox's main bank account. A reconciliation explains why the bank's number and the company's books differ for the same cash, and shows that once every difference is explained, both land on the same true balance. At September 30, the books say 3,016,600.35 dollars and the bank says 2,896,663.23. That's a gap of about 120 thousand dollars to explain. It takes most of a day. Priya wants AI to do the grunt work, and Walt is willing to try, but he wants to be able to check it.

3 Why not just ask AI to reconcile the whole thing in one request? Because a reconciliation has many small steps: line up two lists, find the lines that don't match, decide what kind of difference each one is, adjust the right balance, and add it up. When all of that arrives as one polished answer, a mistake in one step is buried inside it. You can't tell where to look, and AI may fill a gap with something that merely fits. Big jobs fail when the thinking is hidden.

4 First technique: break the job into steps. The name is task decomposition, and when each answer feeds the next request, it's called prompt chaining. For this reconciliation I'd use four stages: match, classify, compute, and challenge. Each stage is its own request, and I check each one before moving on. Why it helps: every step is small enough for me to check, so mistakes get caught early instead of buried. And if stage two looks wrong, I fix it before stage three builds on it. You'll design your own stages in your own words.

5 Second technique: ask for reasoning and assumptions. This is called chain-of-thought prompting, plus stating assumptions. When AI decides that a line is an outstanding check, I ask why, which line it relied on, and what it assumed. This isn't a magic phrase. Many newer AI models reason internally, but I still can't see or check that reasoning unless it's in the answer. Why it helps: a visible reason is something I can check against the source, and a named assumption shows me exactly where a human judgment is needed.

6 Third technique: self-critique. After AI finishes, I ask it to act as a skeptical reviewer of its own work: what might it have missed, counted twice, or put on the wrong side? This is also called red-team prompting. Why it helps: a second pass with a different job, find what's wrong, often catches things the first pass didn't. But here's the catch. AI can look at its own answer and say it's all good when it isn't. So a self-review is a clue, never proof.

7 Fourth technique: self-consistency. I run the same task a second time in a brand-new chat, so it can't see the first answer, and then I compare. Where the two runs agree, I'm a bit more confident. Where they differ, I've found exactly the items to check first. Why it helps: it tells me where to spend my checking time. But agreement isn't proof, because both runs can make the same mistake.

8 Fifth technique, and the one that matters most: verification against the source. I do three things myself. One, recompute the totals in a spreadsheet. Two, trace each reconciling item to a specific line in the files, by date, reference, and amount. Three, look for what's missing or invented. If AI adds an unexplained adjustment just to make the two sides agree, that's called a plug, and it's a red flag. Why it helps: it turns it looks right into I confirmed it.

9 Let me show you what checking looks like, using a made-up example so I don't spoil your assignment. Suppose AI reconciles a small Velox tooling account and says the outstanding checks total 12,900 dollars. I open the source and find three checks: 5,400, 4,300, and 2,900. I add them in a spreadsheet: 12,600. AI said 12,900, so that's a 300 dollar slip, and it would have thrown off the whole reconciliation. Now suppose it also says a miscellaneous adjustment of 300 closes the gap. I search the source for it and find nothing. That's a plug. I'd fix the total, remove the plug, and recheck. Two checks, a few minutes, and I know exactly where the problem was.

10 This connects to area 2 of the Velox Responsible AI Policy: accuracy. If your name is on a number, you're accountable for it, whether you or AI produced it. That means checking AI output against the source, and being honest about how far you checked. In this week's ethics checkpoint you'll review an AI-prepared reconciliation, decide which claims hold up, and decide what you'd do before it goes to Walt. A reconciliation that forces the two sides to agree is worse than one that honestly says something is unresolved.

11 Here's your week. You'll first try the reconciliation in one request, in your own words. Then you'll learn the techniques and review an AI's work for accuracy. Then you'll run your own four-stage chain, check the results against the source files in a spreadsheet, run it a second time, and write your final reconciliation. At the end you'll save your best prompts as a reusable template and send Walt a short note. Write your own prompts. Let's get to work.