Module 6 - Slide Deck Video Transcription

1 Hi, I'm Jordan Reyes from Tallgrass Advisory, a fictional character presented by an AI avatar, on purpose, because this course is about using AI well and being open about it. Welcome to your third rotation: Internal Audit, working for Nadia Ferreira. This week you'll have AI pull facts out of messy records and test a large set of payments for unusual patterns. Then you'll check what it found and report it carefully, because in an audit, a pattern is a lead and never a verdict.

2 Here's Nadia's situation. Last night the ethics hotline took an anonymous tip: something is not right with a vendor that was set up very fast in Purchasing. The caller gave no names. Nadia has the third-quarter payments ledger, about two thousand payments, plus the vendor list, fourteen invoice records exactly as they arrived, and her own notes. She needs to know whether anything needs follow-up. And she told you something important: she decides what happens next, and she will not put anyone's name in a report on the strength of a pattern.

3 First, the auditing idea, and it's short. One common kind of fraud is a fake vendor: a company that exists only on paper, so the payments go to whoever set it up. Auditors look for signs, called indicators. Invoices just under an approval limit, so no second person has to look. Several invoices on one day. No purchase order. Urgent, pay-today requests. A vendor set up recently, or set up and approved by the same person. But here's the rule that matters most: an indicator is not proof. Each one has innocent explanations, and fraud is decided by people who gather documents, not by a spreadsheet.

4 Now the tool that makes this practical. Benford's Law says that in many real sets of money amounts, the first digit is not equally likely. About thirty percent of amounts start with 1, and the share falls steadily to under five percent for 9. People who invent amounts, or keep them just under a limit, usually don't follow that pattern. So you count first digits, and look for digits that are off by more than about two percentage points. But honest data breaks the pattern too. A fixed monthly fee makes every payment start with the same digit. So a gap is a reason to look, not proof.

5 Let me show you with a made-up example, so I don't give away your assignment. Maple Ridge Dairy Co-op is a small Velox supplier. I gave Claude three invoice records and a ledger of 1,106 payments for one quarter. My first request asked Claude to extract the three records into a table, using a fixed set of columns, writing "not stated" instead of guessing, and quoting the words it used for the amount. My second request asked it to use code to

run a Benford test, name the vendors behind any gap, give innocent explanations, and say what it could not tell. And I told it not to say anyone committed fraud.

6 Here is the first thing Claude returned, and this is its real output. Look at the dates and amounts. The first record said only May 14, with no year, so Claude wrote not stated rather than assume. The second was written 05/04/2026, and Claude noted that the order of month and day wasn't clear. The third said about 3.2 thousand dollars, final invoice to follow, so Claude did not turn that into a number. That is what you want: a table you can count, where a missing fact is shown as missing, not quietly filled in.

7 Next, the Benford test. Claude used code, so these counts are exact. Digit 2 is off by plus four and a half points. Digit 6 is off by plus two point three. Every other digit is within two points. Claude also ran an overall test and said the first-digit pattern as a whole doesn't look like Benford. But notice what comes next, because this is where the real work is. A gap tells you where to look. It doesn't tell you why.

8 Claude traced each gap to a single vendor. Digit 2 came from a trucking lease: all sixty payments were exactly 2,450 dollars. Without that vendor, digit 2 was exactly as expected. Claude's likely innocent explanation: a fixed monthly lease for about twenty trucks, to be checked against the contract. Digit 6 came from a dairy equipment vendor: all forty-six payments between about 6,150 and 6,960, all approved by the same person. Its innocent explanations: a narrow price band, or one approver who owns the account. And Claude said what a person would need to see, like purchase orders and receiving records. It ended by writing, nothing here says anyone committed fraud.

9 Now the techniques behind that demo. First, structured extraction. Records arrive as emails, scans, and forms. To count or compare anything, first put the facts in a table. Tell the AI the columns, with a format for each, what to write when a fact is missing, which is not stated, never a guess, and to quote the words it used for the key fact. Practitioners call this structured extraction, or schema-based extraction. Why it helps: it turns text you can't sort into a table you can, the not-stated rule keeps AI from guessing, and the quote lets you spot-check any row in seconds. In your assignment you'll write your own extraction template.

10 Second, let AI count with code. Counting first digits across two thousand rows by reading them is exactly what AI does badly. It can sound sure and be off by dozens. So ask Claude to use code on the file. Practitioners call this tool use, specifically code execution. Code counts exactly, and it can show you what it ran. You still check one or two counts yourself in a spreadsheet. And the third technique you saw in the demo: after a flag, ask for the innocent explanations and what you'd need to see to check them. That's alternative-

explanations prompting, a kind of devil's-advocate prompt. It shows you which flags are probably harmless.

11 When you write your report, start with the most important sentence: what the data show, and that they are indicators for a person to follow up, not proof. Then the few numbers that matter, the innocent explanations you considered, what you don't know, and the next step for a person, like getting a contract or the full bank details. Refer to people only by ID code. Never accuse anyone, and never write that anyone committed fraud. In this case, any of the three next steps can earn full credit, as long as your numbers and reasoning support it.

12 This week's ethics checkpoint is area five of the Velox Responsible AI Policy: human accountability. A person owns every decision and every signature, and AI never signs off. AI can find patterns, draft text, and suggest next steps. But it must not decide that someone is guilty, accuse or name a person in a report, sign or send a report, or quietly drop a finding. A named person reviews and signs the report, decides which flags to keep and why, and decides who sees information about individuals. In Step 4 you'll judge five AI suggestions, and in Part 8 you'll say which person is accountable for which decision.

13 Here's your week. First you'll try the tip in a plain chat. Then you'll learn the ideas and judge five AI suggestions for accountability. Then you'll write your own extraction template, run it on the fourteen Velox records, and have Claude test the payment ledger. You'll check two findings yourself, and write Nadia a one-page report of indicators. You'll also choose a lens, one of six equal options: accounting, finance, marketing, management, supply chain, or information systems. Write your own prompts, check what the AI finds, and let a person decide. Let's get to work.