Module 8 - Slide Deck Video Transcription

1 Hi, I'm Jordan Reyes from Tallgrass Advisory, a fictional character presented by an AI avatar, on purpose, because this course is about using AI well and being open about it. For seven weeks you rotated through Velox departments. This week Tallgrass becomes Velox's AI implementation partner, and you become a builder. The project is the Month-End Machine. You already own the parts. Now you wire them together, and then Walt Brenner, the Controller, tries to break it.

2 Here's Walt's deal. Every month his team does the same close: reconcile the bank account, check the variance report, scan the payments, and write him a memo. Tallgrass says AI can do the first pass. Walt has heard that before. So you build a machine, and he tests it on three months of simulated Velox files. October is practice: you may see what you missed and fix it. November is the Gauntlet. Walt planted problems you have not seen, and you run it once. December is a fresh month, to find out whether your fixes work on something new. A person still signs the memo.

3 First, two words. A workflow is a fixed series of steps. You decide the steps, the order, and where a person must look, and the AI does the work inside each step. An agent is an AI that chooses its own steps. You build agents next week. For a task that must be the same every month and that somebody signs, start with a workflow: it is predictable, and you can test it. Your machine has four stages: reconcile cash, explain the variances, scan the payments, and write the close memo. It stops twice, at Gate 1 and Gate 2, for a person.

4 Each stage reuses a piece you already built. Stage 1 reuses your Module 3 reconciliation chain: match, classify, compute, challenge. Stage 2 reuses your Module 2 variance template and your Module 4 habit that the numbers say what and the notes say why. Stage 3 reuses your Module 6 payment checks and your Module 7 rule to say, not in the sources, which here becomes, not in the data. Stage 4 reuses your prompt constraints. The technique has a name: prompt chaining, also called a multi-step workflow or pipeline. Each stage is small enough to check, and its output feeds the next. And Walt's thresholds live in one rules file, not inside the stages, so he can change a rule without anyone rewriting the machine.

5 Gates. A gate is a point where the machine writes the word gate, stops, and waits for a person. Gate 1 comes after the cash reconciliation, because the adjusted balance feeds the memo, and an error there spoils everything after it. Gate 2 comes after the payment scan, just before the memo, so a person decides what goes in front of the Controller. The

technique is called human-in-the-loop checkpoints. It helps because the AI does the heavy lifting while a person keeps the judgment and the accountability. Two warnings. A gate approved without reading is a rubber stamp, not a control. And no exceptions found is not the same as looked and found nothing.

6 What makes it a machine is the instructions. They live in a Claude Project: the rules file and the vendor list are knowledge files, the instructions say how to work, and each month you attach that month's files to a new chat. Good instructions name the stages and the gates, and require code for every calculation. They also include guardrails. A guardrail is a stop rule. Use only the files. Name every item by its code. Write, not in the data, instead of guessing. If a file is missing, stop and say so. Never imply that anyone committed fraud, and never name individuals. A machine that stops when something is wrong is safer than one that improvises.

7 How do you know it works? Not because the memo sounds confident. You test it on a month where you know the answers. Walt plants problems. You count how many your machine caught, how many it missed, and how many false alarms it raised. Practitioners call caught-out-of-planted recall, and real-out-of-flagged precision. Say a test month has twelve planted problems. The machine lists ten items, and eight are planted: it caught eight of twelve, with two false alarms. A machine that flags everything catches everything and is useless. Then the rule that matters most: fix the rule, not the practice month. If the machine missed a duplicate payment, the fix is a general instruction to compare every payment with every other. A note about one payment passes October and fails December. That is overfitting. After any change, re-run an earlier month to check you broke nothing. That is regression testing.

8 Let me show you with a made-up company, Lakeshore Marine Supply, so I don't give away your answers. We built its machine the way you will build yours. On the first try, we attached only three of the five files. Here is Claude's real answer. I can't start the close yet. Two of the files I work from are missing: the bank statement and the GL cash ledger. My instructions say that when a file is missing I stop and tell you rather than improvise. I haven't run any stage and haven't produced any numbers. That is the guardrail working.

9 After we attached the files, Stage 1 matched the bank statement to the ledger item by item. It listed four routine items, which are not exceptions, and three exceptions: RTN0822, a nine thousand six hundred dollar returned item not in the books; CHK51004, which the bank cleared at six thousand one hundred eighty dollars but the books recorded at six thousand eight hundred ten; and WIRE9025, a four hundred twenty-five dollar fee the bank itself calls charged in error. Claude used Python to calculate both adjusted balances. Both

came to five hundred thirty thousand, four hundred thirty dollars and eighty-three cents. The difference is zero. Then it wrote GATE 1 and waited.

10 Stage 2 is where I want you to notice something. Three lines met the test. For L01, revenue, the note said the boat show season ended early: explained. For L06, marketing, booth deposits were paid in August instead of September: explained. For L02, cost of goods sold, the note said fiberglass resin price up slightly, but cost was three hundred forty-five thousand dollars under budget. Claude said: this does not explain it. A price increase would push cost up. The reason cost is lower is not in the data. It did not accept a note just because a note existed, and it did not guess. That line goes to a person.

11 Here is the memo, after both gates. It opens with the adjusted cash balance and the number of exceptions. Then it lists items by stage, using codes only: three cash items, the three variance lines, and four payment problems, a duplicate, a possible split, a payment over the approval limit with no second approval, and a vendor not on the approved list. It names nobody, says not in the data where the files are silent, and ends with a line for a person to sign. For this demonstration we told the machine to treat both gates as approved. You will not do that. You will stop and read at each gate.

12 The ethics checkpoint is human accountability, policy area five. AI may do the first pass: match, check, calculate, draft. A named person, a role like the Controller or an accountant on the close team, reviews the work and answers for it, and a machine can never be that person. So the memo ends with a signature line. You will judge five suggestions, such as approving a gate without reading, or emailing the memo automatically with the machine listed as reviewer, and you will decide which ones keep a person accountable.

13 Here's your week. Try one plain chat with the October files and no rules, and see what Claude decides on its own. Learn the ideas and judge five suggestions. Map your saved pieces to the four stages, write your instructions, and build your Project. Rehearse on October and score it. Take Walt's Gauntlet on November, check two findings yourself, write a one-page card for your machine, and choose a lens for your final-project idea. Then get my feedback, fix your instructions with general rules, and test on December. Quiz and reflect. Walt is waiting.